

✓ Import Lib

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline

import warnings
warnings.filterwarnings('ignore')

# importing Stopwords
import nltk
from nltk.corpus import stopwords
import string

# models
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.naive_bayes import MultinomialNB
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC

# train test split
from sklearn.model_selection import train_test_split, GridSearchCV

# Pipeline
from sklearn.pipeline import Pipeline

# score
from sklearn.metrics import confusion_matrix, classification_report, ConfusionMatrixDisplay
from sklearn.metrics import accuracy_score
```

✓ Import Data

```
df=pd.read_csv('/kaggle/input/spam-ham-dataset/spam_dataset.csv')
```

✓ Basic Analysis

```
df.head()
```

	Unnamed: 0	label	text	label_num
0	605	ham	Subject: enron methanol ; meter # : 988291\r\n...	0
1	2349	ham	Subject: hpl nom for january 9 , 2001\r\n(see...	0
2	3624	ham	Subject: neon retreat\r\nho ho ho , we ' re ar...	0
3	4685	spam	Subject: photoshop , windows , office . cheap ...	1
4	2030	ham	Subject: re : indian springs\r\nthis deal is t...	0

```
df.describe()
```

	Unnamed: 0	label_num
count	5171.000000	5171.000000
mean	2585.000000	0.289886
std	1492.883452	0.453753
min	0.000000	0.000000
25%	1292.500000	0.000000
50%	2585.000000	0.000000
75%	3877.500000	1.000000
max	5170.000000	1.000000

```
df['label'].value_counts()
```

```
label
ham      3672
spam     1499
Name: count, dtype: int64
```

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 5171 entries, 0 to 5170
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Unnamed: 0    5171 non-null   int64
1   label         5171 non-null   object
2   text          5171 non-null   object
3   label_num     5171 non-null   int64
dtypes: int64(2), object(2)
memory usage: 161.7+ KB
```

```
# adding new column as length of the text
```

```
df['length'] = df['text'].apply(len)
```

```
df.head()
```

	Unnamed: 0	label	text	label_num	length
0	605	ham	Subject: enron methanol ; meter # : 988291\r\n...	0	327
1	2349	ham	Subject: hpl nom for january 9 , 2001\r\n(see...	0	97
2	3624	ham	Subject: neon retreat\r\nho ho ho , we ' re ar...	0	2524
3	4685	spam	Subject: photoshop , windows , office . cheap ...	1	414
4	2030	ham	Subject: re : indian springs\r\nthis deal is t...	0	336

EDA

```
# plot for count of spam and ham in data
```

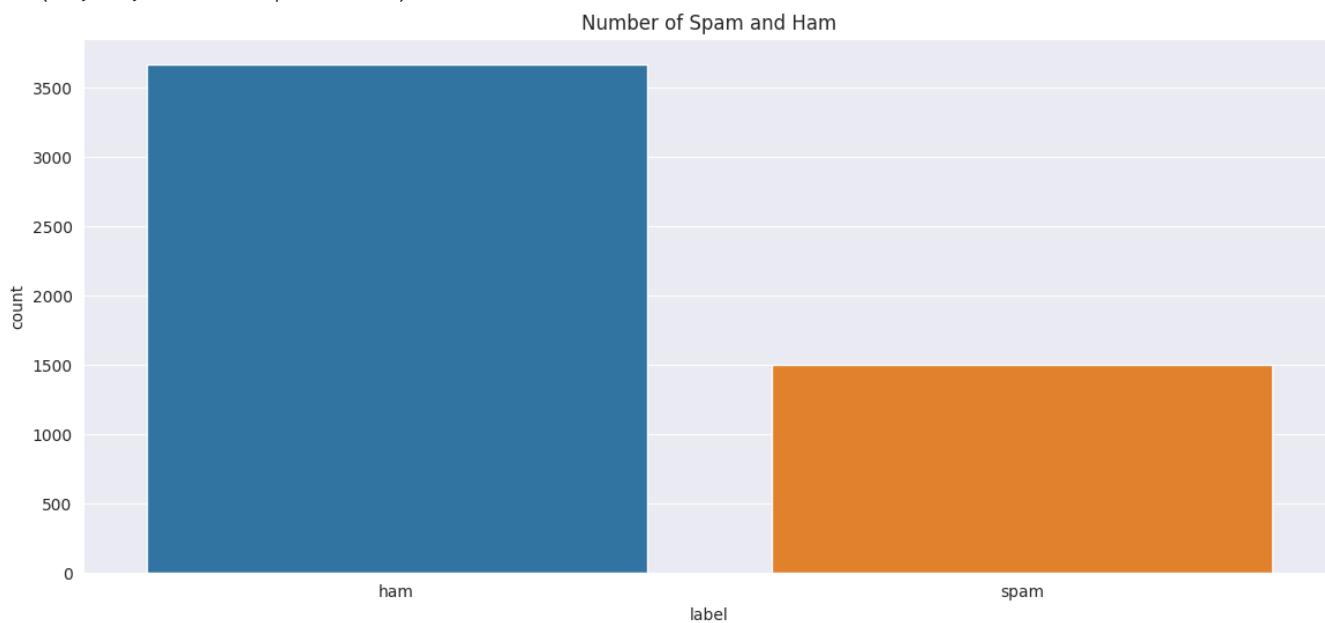
```
plt.figure(figsize=(14,6))
```

```
sns.set_style('darkgrid')
```

```
sns.countplot(x='label',data=df)
```

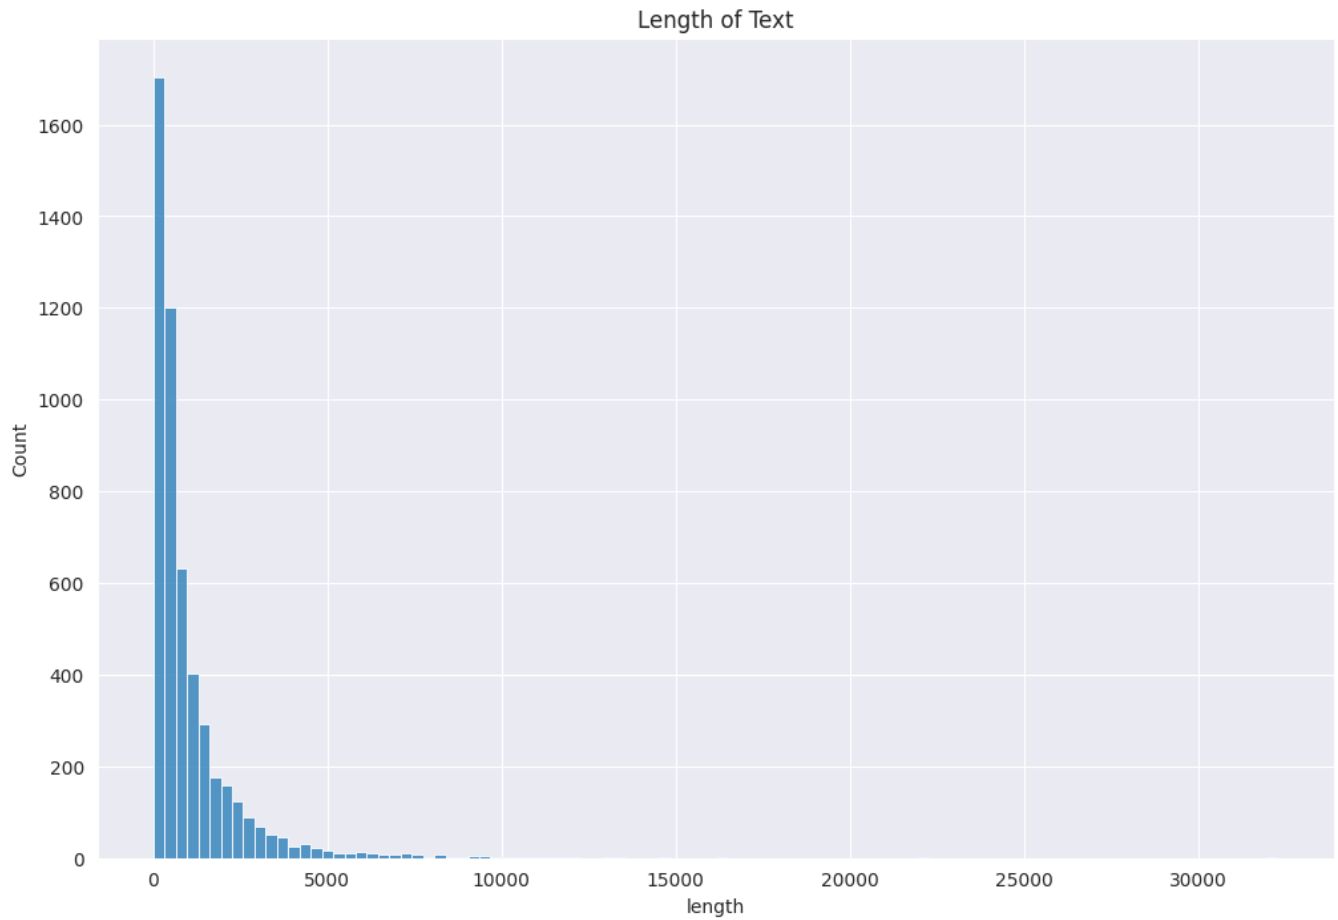
```
plt.title('Number of Spam and Ham')
```

Text(0.5, 1.0, 'Number of Spam and Ham')



```
# Plot for distribution lenth of text
plt.figure(figsize=(12,8))
sns.histplot(x='length',data=df,bins=100)
plt.title('Length of Text')
```

```
Text(0.5, 1.0, 'Length of Text')
```




```
# maximum lenth text
df[df['length']==df['length'].max()]['text']

949    Subject: fw : " red , white and blue out "\r\n...
Name: text, dtype: object
```

```
# distribution of spam and ham by length of text
df.hist(column='length',by='label',figsize=(12,8))
```

```
array([<Axes: title={'center': 'ham'}>, <Axes: title={'center': 'spam'}>],
      dtype=object)
```



Feature Engineering

```
# function to remove punctuation and stopwords
```

```
def text_process(text):
    non_punc = [char for char in text if char not in string.punctuation]
    non_punc=''.join(non_punc)
    return [word for word in non_punc.split() if word not in stopwords.words('english')]
```

Train Test Split

```
# define X(features),y(target)
```

```
X= df['text']
y=df['label']
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random_state=42)
```

Modles

```
# creating a pipline to model the data
```

```
# pipeline for MultinomialNB
```

```
pipe_mnb = Pipeline([
    ('bow',CountVectorizer(analyzer=text_process)),
    ('tf',TfidfTransformer()),
    ('classifier',MultinomialNB())
])
```

```
# pipeline for Random Forest Classifier
```

```
pipe_rf =Pipeline([
    ('bow',CountVectorizer(analyzer=text_process)),
    ('tf',TfidfTransformer()),
    ('classifier',RandomForestClassifier())
])
```

```
# pipeline for Random Forest Classifier
```


```
pipe_svc =Pipeline([
    ('bow',CountVectorizer(analyzer=text_process)),
    ('tf',TfidfTransformer()),
    ('classifier',SVC())
])
```

```
# fit the data
```

```
pipe_mnb.fit(X_train,y_train)
```

```
pipe_rf.fit(X_train,y_train)
```

```
pipe_svc.fit(X_train,y_train)
```



```
# predict the target feature
```

```
pred_mnb = pipe_mnb.predict(X_test)
```

```
pred_rf = pipe_rf.predict(X_test)
```

```
pred_svc = pipe_svc.predict(X_test)
```

✓ Prediction Accuracy

```
print('The accuracy for Multinomial Classifier:',accuracy_score(y_test,pred_mnb)*100)
print('The accuracy for Random_forest Classifier:',accuracy_score(y_test,pred_rf)*100)
print('The accuracy for SVC:',accuracy_score(y_test,pred_svc)*100)
```

```
The accuracy for Multinomial Classifier: 91.73989455184535
The accuracy for Random_forest Classifier: 97.4223784417106
The accuracy for SVC: 98.82835383714118
```

The SVC predicts better than Random Forest Model and Multinomial.

```
# print confusion matrix and classification report
print ('Classification report on SVC:')
print('\n')
print(classification_report(y_test,pred_svc))
```


Classification report on SVC:

	precision	recall	f1-score	support
ham	1.00	0.99	0.99	1246
spam	0.97	0.99	0.98	461
accuracy			0.99	1707
macro avg	0.98	0.99	0.99	1707
weighted avg	0.99	0.99	0.99	1707

```
# Display confusion matrix for SVC
```

```
sns.set_style('ticks')
ConfusionMatrixDisplay(confusion_matrix(y_test,pred_svc)).plot()
plt.title("Confusion Matrix for SVC")
```

```
Text(0.5, 1.0, 'Confusion Matrix for SVC')
```



```
from sklearn.model_selection import cross_val_score

# Number of folds
k = 5

# Initialize the SVC model in the pipeline
pipe_svc.set_params(classifier=SVC())

# Perform k-fold cross-validation
cv_scores = cross_val_score(pipe_svc, X, y, cv=k)

# Output the results
print(f'CV Scores for each fold: {cv_scores}')
print(f'Average CV Score: {np.mean(cv_scores)}')
```

CV Scores for each fold: [0.98937198 0.98742747 0.99323017 0.98549323 0.99129594]
Average CV Score: 0.9893637578373934